



[< Back to Machine Learning Engineer Nanodegree](#)


Finding Donors for CharityML

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

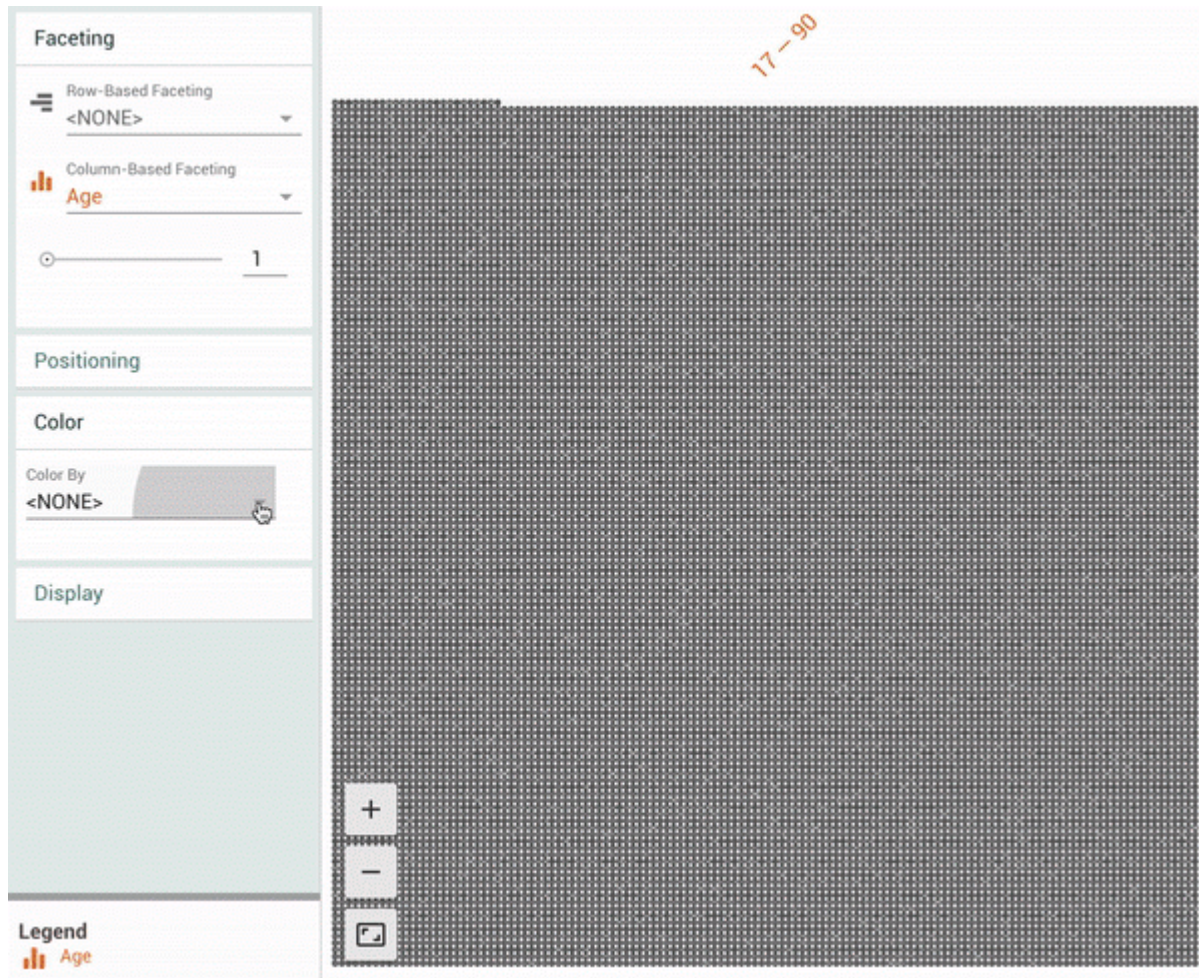


Congratulations on passing your exam!

Bonus

A partnership with Google has released an [open source software](#) to visualize machine learning datasets. And the best part is that they are using the same dataset of this project as live demo! If you want to perform some data

visualization with the data of this project, go to the link <https://pair-code.github.io/facets/> and play around. Here's an example from Google's blog:



Exploring the Data

Student's implementation correctly calculates the following:

- Number of records
- Number of individuals with income >\$50,000
- Number of individuals with income <=\$50,000
- Percentage of individuals with income > \$50,000

Bonus

We can also get some initial statistics with the pandas method [describe](#):

```
data.describe()
```

In addition, a very cool tool for data visualization and correlations is the pairplot of seaborn (install using `pip install seaborn`):

```
In [9]: import seaborn as sns
sns.pairplot(data, palette='Set1', hue='income')
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1169b6050>
```



Preparing the Data

Student correctly implements one-hot encoding for the feature and income data.

Awesome

Well done using the lambda function!

Comment

[This reference](#) provides 7 different encoding strategies. [Binary encoding](#) is a great choice for cases where the number of categories for a given feature is very high.

Evaluating Model Performance

Student correctly calculates the benchmark score of the naive predictor for both accuracy and F1 scores.

Great job calculating the accuracy and the F-score for a Naive predictor!

Comment

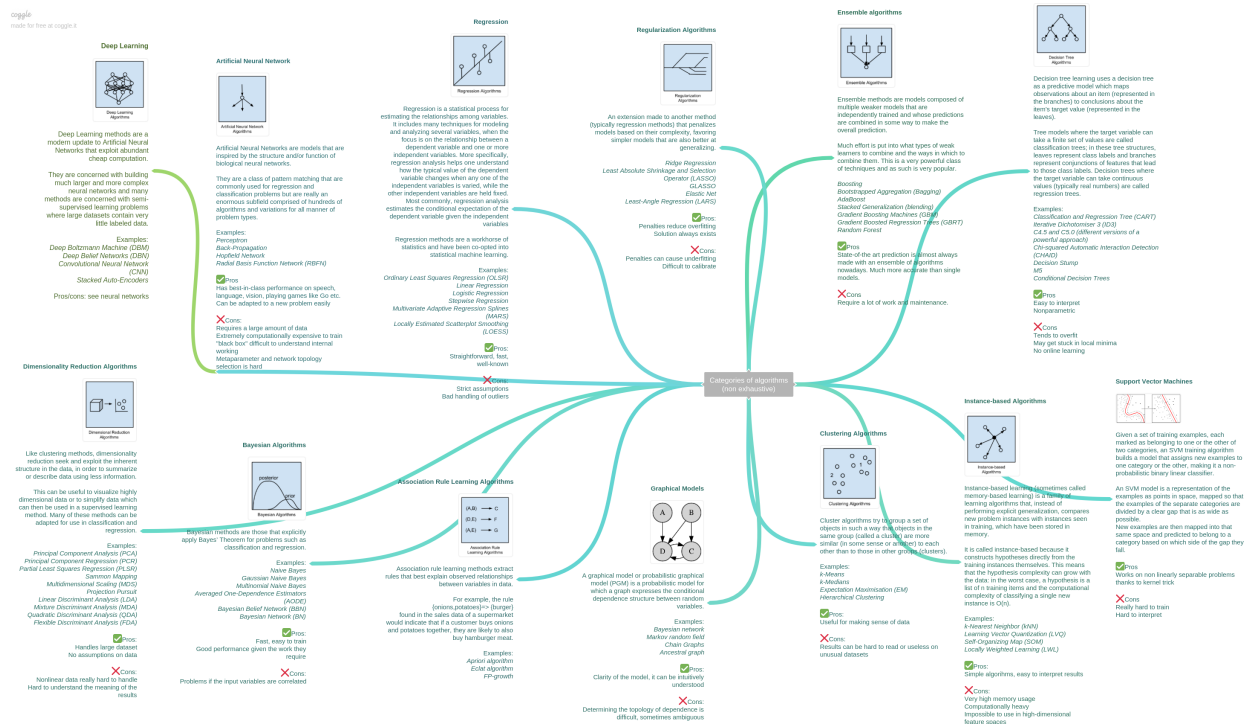
Note that the F-score is higher than the accuracy which seems counter-intuitive since the F-score is a more elaborate calculation. That happens because a value of $\beta = 0.5$ attenuates the influence of false negatives. In other words, this value of β weights more the positive predictions ($>50K$) than the negative one ($\leq 50K$).

The pros and cons or application for each model is provided with reasonable justification why each model was chosen to be explored.

Please list all the references you use while listing out your pros and cons.

Bonus

You can also refer to [this reference](#) when choosing your estimator:



- [This reference](#) provides many visualization resources of machine learning algorithms. I definitely recommend it to get a visual intuition of ML techniques. My favorite is the article [Gradient Boosting explained](#).
- [This reference from sklearn](#) compares the decision boundaries of the major classifiers.
- [This map from sklearn](#) might be also helpful in order to give guidance when choosing an estimator.
- [XGBoost](#) and [Random Forest](#) are two very popular algorithms in data science contests. [This guide](#) might be useful for tuning XGBoost.
- Three suggestion of books in case you are interested in diving deeper into the characteristics of each estimator:
 - [The Elements of Statistical Learning](#)
 - [Bishop: Pattern Recognition and Machine Learning](#)
 - [Machine Learning - Tom Mitchell](#)

Student successfully implements a pipeline in code that will train and predict on the supervised learning algorithm given.

Awesome

Everything looks great here!

Student correctly implements three supervised learning models and produces a performance visualization.

Improving Results


Justification is provided for which model appears to be the best to use given computational cost, model performance, and the characteristics of the data.

I agree with your choice of Gradient Boosting here! It is one of the best estimators for this project and in the analysis is the one which is leading to the highest test score. It is worth noting that it might not be fair to compare different models with their default parameters. For example, Decision Trees tend to not perform well with their default parameters while their tuned model can surpass the tuned SVM.


Student is able to clearly and concisely describe how the optimal model works in layman's terms to someone who is not familiar with machine learning nor has a technical background.

Bonus

This article called [Gradient Boosting explained](#) is great for visualizing GB in action:



In addition, this [video](#) shows AdaBoost (GB's cousin) in action. It might be useful to get an intuition of how this estimator works. I suggest watching it in slow motion:



The final model chosen is correctly tuned using grid search with at least one parameter using at least three settings. If the model does not need any parameter tuning it is explicitly stated with reasonable justification.

Bonus

Although Gradient Boosting usually give better results, AdaBoost has as advantage that we can tune its base estimator using the prefix `base_estimator__`. Here's an example:

```
# TODO: Initialize the classifier
clf = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(random_state=42))

# TODO: Create the parameters list you wish to tune, using a dictionary if needed.
parameters = {'n_estimators':[50, 75, 120],
              'learning_rate':[0.1, 0.5, 1.],
              'base_estimator__min_samples_split' : np.arange(2, 8, 2),
              'base_estimator__max_depth' : np.arange(1, 4, 1)
              }
```

Student reports the accuracy and F1 score of the optimized, unoptimized, models correctly in the table provided. Student compares the final model results to previous results obtained.

Awesome improvement in the model using grid search! Adaboost and Gradient Boosting are usually the estimators with highest scores for this project.

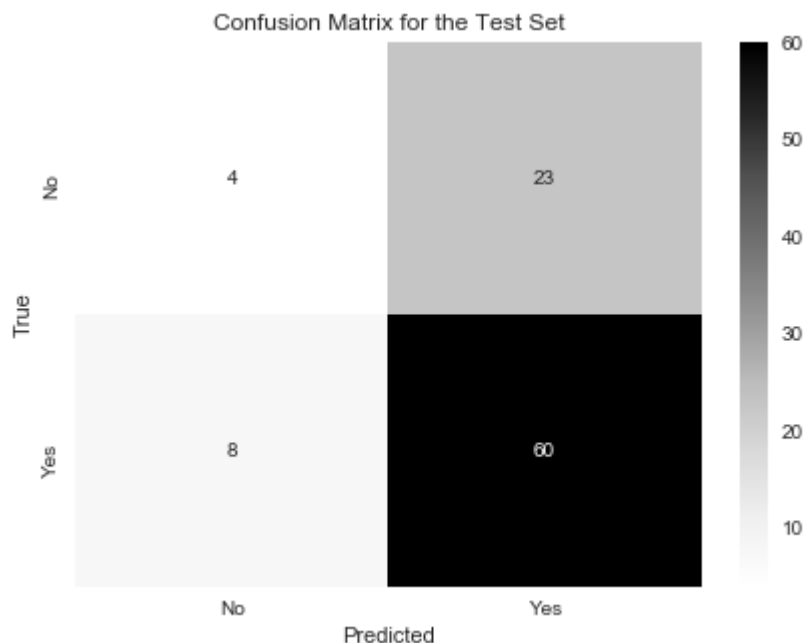
Bonus

You can also check your results with a [Confusion Matrix](#):

```
import seaborn as sns # Install using 'pip install seaborn'
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
%matplotlib inline

cm_test = confusion_matrix(y_test, best_clf.predict(X_test))

plt.figure(figsize=(7,5))
sns.heatmap(cm_test, annot=True, cmap='Greys', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.title('Confusion Matrix for the Test Set')
plt.ylabel('True')
plt.xlabel('Predicted')
```



Bonus 2

In case you are interested in even higher scores, give a try on XGBoost (install using `pip install xgboost`). It works just like any Sklearn estimator. Here's an example to get scores above 0.75:

```
from xgboost import XGBClassifier # pip install xgboost
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer

# Initialize the classifier
clf = XGBClassifier(random_state=0)

# Create the parameters list you wish to tune, using a dictionary if needed.
# HINT: parameters = {'parameter_1': [value1, value2], 'parameter_2': [value1, value
2]}
parameters = { 'max_depth': [2, 3, 5],
                'learning_rate': [0.01, 0.1, 0.2],
                'n_estimators': [50, 100, 150] }

# Make an fbeta_score scoring object using make_scorer()
scorer = make_scorer(fbeta_score, beta=0.5)

# Perform grid search on the classifier using 'scorer' as the scoring method using G
ridSearchCV()
grid_obj = GridSearchCV(clf, parameters, scoring=scorer, verbose=10)

# Fit the grid search object to the training data and find the optimal parameters us
ing fit()
```



```
grid_fit = grid_obj.fit(X_train, y_train)

# Get the estimator
best_clf = grid_fit.best_estimator_

print("Best parameters: ")

# Make predictions using the unoptimized and model
predictions = (clf.fit(X_train, y_train)).predict(X_test)
best_predictions = best_clf.predict(X_test)

# Report the before-and-afterscores
print("Unoptimized model\n-----")
print("Accuracy score on testing data: {:.4f}".format(accuracy_score(y_test, predictions)))
print("F-score on testing data: {:.4f}".format(fbeta_score(y_test, predictions, beta = 0.5)))
print("\nOptimized Model\n-----")
print("Final accuracy score on the testing data: {:.4f}".format(accuracy_score(y_test, best_predictions)))
print("Final F-score on the testing data: {:.4f}".format(fbeta_score(y_test, best_predictions, beta = 0.5)))
```

With more tuning, we can get scores higher than 0.78.

Feature Importance

Student ranks five features which they believe to be the most relevant for predicting an individual's income. Discussion is provided for why these features were chosen.

Student correctly implements a supervised learning model that makes use of the `feature_importances_` attribute. Additionally, student discusses the differences or similarities between the features they considered relevant and the reported relevant features.

Comment

It is worth noting that each model with `feature_importances_` might return different top predictive features depending on their internal algorithm implementation.

Suggestion

You can also use the attribute `feature_importances_` from `best_clf` since it is already tuned so you will have a better choice of top 5 features:

```
importances = best_clf.feature_importances_
```

Bonus

Besides the feature importance, an alternative is to use calculate the significance of each feature using the weights of a Logistic Regression (LR) model. The significance is calculated dividing the coefficients of the LR by their standard errors (also know as Z-score). You can check more details in the Chapter 4.4.2 of [The Elements of Statistical Learning](#). According to the reference, a Z-score higher than 2 is significant at the 5% level.

Student analyzes the final model's performance when only the top 5 features are used and compares this performance to the optimized model from Question 5.

Comment

An alternative strategy for reducing the number of features is to use dimensionality reduction techniques (PCA for example). Then, we could pick only the top descriptive features for training the model. You will see more details about PCA in the next module.

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)